

Domain-Specific Acceleration and Pragmatic Auto-Parallelization of Legacy Scientific Code in FORTRAN 77 using Source-to-Source Compilation

Wim Vanderbauwhede, Gavin Davidson

School of Computing Science, University of Glasgow, UK



- ▶ A lot of scientific code is “legacy” code written in FORTRAN 77
- ▶ Massively parallel accelerators such as GPGPUs, manycores and FPGAs are powerful and affordable tools for scientists to speed up simulations of complex systems.
- ▶ Porting code to such devices requires a detailed understanding of heterogeneous programming tools and effective strategies for parallelization.
- ▶ We present our work on source to source compilation to
 - ▶ transform sequential FORTRAN 77 legacy code
 - ▶ into OpenCL-accelerated programs
 - ▶ with auto-parallelized kernels
 - ▶ without need for directives or extra information from the user.



Accelerator Programming

- ▶ **OpenCL**: open standard for heterogeneous computing
- ▶ Main advantage of OpenCL over proprietary solutions such as e.g. CUDA (to which it is very similar): it is supported by a wide range of devices, including multicore CPUs, manycores, GPGPUs and FPGAs.
- ▶ Programmer writes
 - ▶ one or more kernels that are run directly by the accelerator and
 - ▶ a host program that is run on the system's main CPU.
- ▶ The host program handles memory transfers to the device and initializing computations
- ▶ The kernels perform the parallel computations.



- ▶ From the programmer perspective, OpenCL is
 - ▶ very flexible
 - ▶ quite low level
 - ▶ requires a lot of complex code to be written
- ▶ Considerable barrier for adoption
- ▶ No official Fortran support for OpenCL:
 - ▶ Host API is C/C++
 - ▶ Kernel language is based on a subset of C99.

- ▶ A compiler to
 - ▶ transform FORTRAN 77 code into accelerator-ready Fortran 95
<https://github.com/wimvanderbauwhede/RefactorF4ACC>
 - ▶ transform the Fortran program into a Fortran-OpenCL program with automatic parallelisation <https://github.com/wimvanderbauwhede/AutoParallel-Fortran>
 - ▶ compile Fortran 95 kernels into the OpenCL C language
<https://github.com/wimvanderbauwhede/RefactorF4ACC>
- ▶ A Fortran API for OpenCL
<https://github.com/wimvanderbauwhede/OpenCLIntegration>

Source-to-source compilation

- ▶ A conventional compiler consumes source code and produces binaries.
- ▶ A source-to-source compiler produces a transformed version of the original source.
 - ▶ e.g. refactoring, parallelization or translation to a different language.
- ▶ The advantage of this approach is that the resulting code can be modified by the programmer if desired and compiled with a compiler of choice.
- ▶ We combine source-to-source compilation with whole-program analysis:
 - ▶ most compilers only consider a single source file as the context.
 - ▶ our compiler analyses the complete source of a program and establishes the relationships between all entities



Maintainable and Extensible

In a first step the compiler transforms the FORTRAN 77 program into modern, maintainable, extensible and accelerator-ready Fortran 95 code

- ▶ No IMPLICIT typing: our compiler creates explicit type declarations for all variables.
- ▶ Fully explicit INTENT: our compiler infers the INTENT for all subroutine and function arguments.
- ▶ Modules with export lists: our compiler converts all non-program code units into modules USEd with an explicit export (ONLY) declaration.



Accelerator-ready

- ▶ Most current accelerators have a separate memory space from the host memory.
- ▶ It is therefore crucial to separate the memory spaces of the kernel and the host programs.
- ▶ Our compiler transforms global variables (e.g COMMON block variables) into subroutine arguments across the complete call tree of the program.

OpenCL-accelerated Code

1. *Allow any subroutine in the code to be offloaded* to an accelerator using whole-source refactoring.
2. *Minimization of the data transfer* between the host and the accelerator by eliminating redundant transfers.
Includes determining which transfers need to be made only once in the run of the program.
3. *Pragmatic auto-parallelization* of the code to be offloaded to the accelerator by identification of parallelizable maps and folds (reductions).
Includes partial parallelization of loops and fusion of loops with non-identical bounds.



Code Transformation Validation

- ▶ We have validated the code transformation performance of the compiler on the *NIST FORTRAN78* test suite (72,4K loc) as well as on following real-world codes:
 - ▶ *Large Eddy Simulator for Urban Flows* (LES), a high-resolution turbulent flow model (1.4K loc)
 - ▶ Shallow water component of the *Gmodel ocean model* (1.5K loc)
 - ▶ *Flexpart-WRF*, a particle dispersion simulator (13.8K loc)
 - ▶ *Linear Baroclinic Model*, an atmospheric climate model (39.3K loc)
- ▶ All these codes are successfully refactored:
 - ▶ no errors on code generation
 - ▶ no errors on compilation
 - ▶ transformed code performance and results are identical to original code



Automatic Parallelization Validation

- ▶ Test case: Large Eddy Simulator for Urban Flows (LES)
- ▶ Developed by the Disaster Prevention Research Institute of Kyoto University and the Japan Atomic Energy Agency:
 - ▶ Generates turbulent flows by using mesoscale meteorological simulations.
 - ▶ Explicitly represents the urban surface geometry.
 - ▶ Used to conduct building-resolving large-eddy simulations of boundary-layer flows over urban areas under realistic meteorological conditions.
 - ▶ Essentially solves the Poisson equation for the pressure, using Successive Over-Relaxation
 - ▶ Written in Fortran-77, single-threaded, about a thousand lines of code.

LES Code Structure – Functional

- LES structure is sequential in each time step:

- velnw:** Update velocity for current time step
- bondv1:** Calculate boundary conditions (initial wind profile, inflow, outflow)
- velfg:** Calculate the body force
- feedbf:** Calculation of building effects (Goldstein damping model)
- les:** Calculation of viscosity terms (Smagorinsky model)
- adam:** Adams-Bashforth time integration
- press:** Solving of Poisson equation using SOR

```

program main
c-----
include 'common.sn' ! Parameter & variable setting, declaration

call set
call grid

call timdata
call init
call ifdata

c--main loop
do 1000 n = n0,nmax
    time = float(n-1)*dt
c-----calculate turbulent flow-----c
    call velnw
    call bondv1
    call velfg

    if(iibf.eq.1) then
        call feedbf !calculate building effect
    end if

    call les
    call adam
    call press
c-----data output-----c
    call timeris
    call aveflow
    if(ianime.eq.1) then
        call anime
    endif

c
    if(n.eq.nmax) then
        stop
    end if
1000 continue
c
end program
    
```



OpenCL LES Evaluation

▶ Platform

- ▶ Host platform: Intel Xeon CPU E5-2620 0 @ 2.00GHz, 6-core CPU with hyperthreading (12 threads), AVX, 32GB RAM, cache 15MB
- ▶ GPU platform: NVIDIA GeForce GTX TITAN, 980 MHz, 15 compute units, 16GB RAM, OpenCL 1.1 CUDA 6.5.14

▶ Original LES code on CPU (reference)

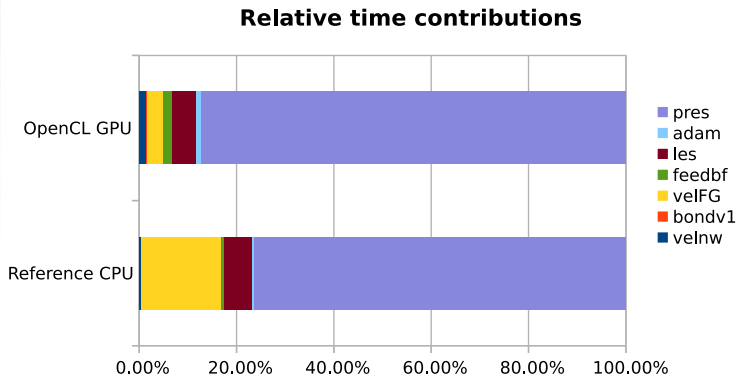
- ▶ Compiler: gfortran 4.8.2, flags for auto-vectorization and auto-parallelization:
 - ▶ `-Ofast -floop-parallelize-all -ftree-parallelize-loops=12 -fopenmp -pthread`
 - ▶ Auto-parallelization provides only 4% speed-up because the most time-consuming loops are not parallelised

▶ Our compiler:

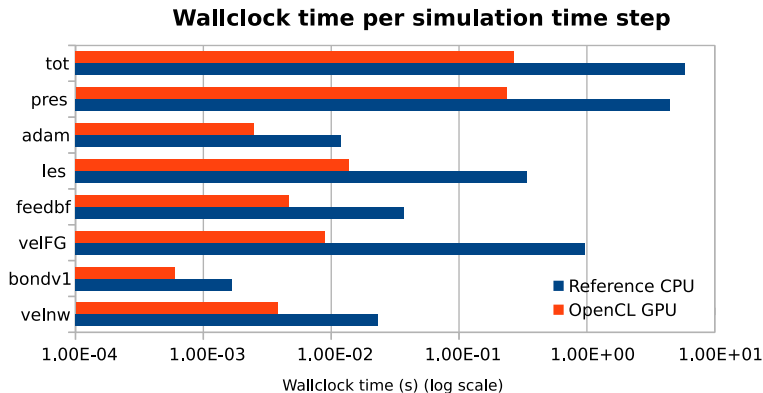
- ▶ Auto-parallelizes all 34 loop nests in the code base
- ▶ Produces a complete OpenCL-enabled code base.
- ▶ Runs on GPU and CPU
- ▶ 20x speed-up compared to the original code on CPU



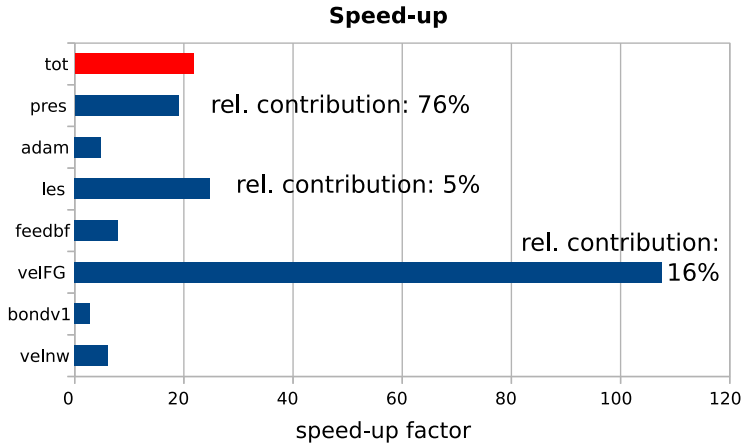
Breakdown per subroutine (domain size 300x300x90, SOR 50 iterations)



Wallclock time comparison



Speed-up



Conclusion

- ▶ We have developed a proof of concept compiler for
 - ▶ OpenCL acceleration and pragmatic auto-parallelization
 - ▶ of domain-specific legacy FORTRAN 77 scientific code
 - ▶ using whole-program analysis and source-to-source compilation.
- ▶ Future work
 - ▶ improving the compiler to extract more parallelism from the original code and improve the performance
 - ▶ FPGA back-end
 - ▶ CUDA and OpenMP back-ends

The authors acknowledge the support of the EPSRC (EP/L00058X/1).



Thank You!

